

## Using tidymodels

Let's first split our data into training and testing datasets:

```
set.seed(1)
split <- initial_split(data = gss_subset, prop = 3/4)
gss_train <- training(split)
gss_test <- testing(split)
```

Next, let's use 5-fold cross validation:

```
folds <- rsample::vfold_cv(gss_train, v = 5)
```

Now, let's make our `recipe()` and `workflow()` that will be used for each of our models:

```
# Create the recipe
gss_recipe <- recipe(partyid ~ ., data = gss_train) %>%
  step_rm(year) %>%
  step_dummy(all_nominal(), -all_outcomes()) %>%
  step_zv(all_predictors())

# Create the workflow
gss_workflow <- workflow() %>%
  add_recipe(gss_recipe)

# View the workflow
gss_workflow
```

```
## == Workflow =====
## Preprocessor: Recipe
## Model: None
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_rm()
## * step_dummy()
## * step_zv()
```

Now, we can begin to specify our models for our model stack:

```
# Basic Logistic regression specification
basic_logreg_spec <- logistic_reg() %>%
  set_engine("glm")

# Add logistic regression to workflow
basic_logreg_workflow <- gss_workflow %>%
  add_model(basic_logreg_spec)

# Cross validation
set.seed(13)
basic_logreg_resamples <- fit_resamples(
```

```

basic_logreg_workflow,
resamples = folds,
control = control_stack_resamples()
)

```

## The workflow being saved contains a recipe, which is 0.4 Mb in memory. If this was not intentional, p

```

# Penalized Logistic regression specification
logreg_spec <- logistic_reg(penalty = tune(),
                           mixture = tune()) %>%
  set_engine("glmnet")

# add grid
lr_reg_grid <- tidyr::crossing(
  penalty = 10 ^ seq(-6, -1, length.out = 20),
  mixture = c(0.05, 0.2, 0.4, 0.6, 0.8, 1)
)

# Add logistic regression to workflow
logreg_workflow <- gss_workflow %>%
  add_model(logreg_spec)

# Tuning hyperparameters
set.seed(13)
logreg_resamples <- tune_grid(
  logreg_workflow,
  resamples = folds,
  grid = lr_reg_grid,
  control = control_stack_grid()
)

```

## The workflow being saved contains a recipe, which is 0.4 Mb in memory. If this was not intentional, p

```

# LDA specification
lda_spec <- discrim_linear(penalty = tune()) %>%
  set_engine("mda")

# Add LDA to workflow
lda_workflow <- gss_workflow %>%
  add_model(lda_spec)

# Fit with our cross validation
set.seed(13)
lda_resamples <- tune_grid(
  lda_workflow,
  resamples = folds,
  control = control_stack_grid()
)

```

## The workflow being saved contains a recipe, which is 0.4 Mb in memory. If this was not intentional, p

```

# Specify random forest
rf_spec <- rand_forest(mtry = tune(),
                      min_n = tune(),
                      trees = 1000) %>%
  set_mode("classification") %>%
  set_engine("ranger")

# Workflow
rf_workflow <-
  gss_workflow %>%
  add_model(rf_spec)

# tuning
set.seed(13)
rf_res <-
  tune_grid(
    rf_workflow,
    resamples = folds,
    grid = 3,
    control = control_stack_grid()
  )

```

## i Creating pre-processing data to finalize unknown parameter: mtry

## The workflow being saved contains a recipe, which is 0.4 Mb in memory. If this was not intentional, p

```

knn_spec <- nearest_neighbor(neighbors = tune()) %>%
  set_mode("classification") %>%
  set_engine("kknn")

knn_wf <- gss_workflow %>%
  add_model(knn_spec)

# tuning
set.seed(13)
knn_res <-
  tune_grid(
    knn_wf,
    resamples = folds,
    control = control_stack_grid()
  )

```

## The workflow being saved contains a recipe, which is 0.4 Mb in memory. If this was not intentional, p

Now, we can stack these models:

```

gss_stack <- stacks() %>%
  add_candidates(basic_logreg_resamples) %>%
  add_candidates(logreg_resamples) %>%
  add_candidates(lda_resamples) %>%
  add_candidates(rf_res) %>%
  add_candidates(knn_res)

```

```
## Warning: Predictions from the candidates c(".pred_dem_logreg_resamples_1_002", ".pred_dem_logreg_res_
## ".pred_dem_logreg_resamples_1_044", ".pred_dem_logreg_resamples_1_064", ".pred_dem_logreg_resamples_
## ".pred_dem_logreg_resamples_1_086", ".pred_dem_logreg_resamples_1_106", ".pred_dem_logreg_resamples_
## ".pred_dem_logreg_resamples_1_009", ".pred_dem_logreg_resamples_1_029", ".pred_dem_logreg_resamples_
## ".pred_rep_logreg_resamples_1_022", ".pred_rep_logreg_resamples_1_042", ".pred_rep_logreg_resamples_
## ".pred_rep_logreg_resamples_1_064", ".pred_rep_logreg_resamples_1_084", ".pred_rep_logreg_resamples_
## ".pred_rep_logreg_resamples_1_106", ".pred_rep_logreg_resamples_1_007", ".pred_rep_logreg_resamples_
## ".pred_rep_logreg_resamples_1_029", ".pred_rep_logreg_resamples_1_049", ".pred_rep_logreg_resamples_
```

```
gss_stack
```

```
## # A data stack with 5 model definitions and 88 candidate members:
## #   basic_logreg_resamples: 1 model configuration
## #   logreg_resamples: 65 model configurations
## #   lda_resamples: 10 model configurations
## #   rf_res: 3 model configurations
## #   knn_res: 9 model configurations
## # Outcome: partyid (factor)
```

```
gss_stack <- gss_stack %>%
  blend_predictions()
gss_stack
```

```
## -- A stacked ensemble model -----
```

```
##
## Out of 88 possible candidate members, the ensemble retained 11.
## Lasso penalty: 0.01.
```

```
##
## The 10 highest weighted member classes are:
```

```
## # A tibble: 10 x 3
##   member                                type      weight
##   <chr>                                <chr>      <dbl>
## 1 .pred_rep_logreg_resamples_1_100 logistic_reg 1.57
## 2 .pred_rep_logreg_resamples_1_097 logistic_reg 1.36
## 3 .pred_rep_logreg_resamples_1_073 logistic_reg 0.956
## 4 .pred_rep_rf_res_1_2                rand_forest 0.911
## 5 .pred_rep_rf_res_1_3                rand_forest 0.701
## 6 .pred_rep_logreg_resamples_1_112 logistic_reg 0.417
## 7 .pred_rep_lda_resamples_1_07        discrim_linear 0.272
## 8 .pred_rep_logreg_resamples_1_118 logistic_reg 0.203
## 9 .pred_rep_logreg_resamples_1_015 logistic_reg 0.0441
## 10 .pred_rep_lda_resamples_1_04       discrim_linear 0.00112
```

```
##
## Members have not yet been fitted with 'fit_members()'.
```

```
gss_stack <- gss_stack %>%
  fit_members()
gss_stack
```

```
## -- A stacked ensemble model -----
```

```
##
## Out of 88 possible candidate members, the ensemble retained 11.
## Lasso penalty: 0.01.
```

```
##
## The 10 highest weighted member classes are:
```

```
## # A tibble: 10 x 3
##   member                                type      weight
##   <chr>                                <chr>      <dbl>
## 1 .pred_rep_logreg_resamples_1_100 logistic_reg 1.57
## 2 .pred_rep_logreg_resamples_1_097 logistic_reg 1.36
## 3 .pred_rep_logreg_resamples_1_073 logistic_reg 0.956
## 4 .pred_rep_rf_res_1_2                rand_forest 0.911
## 5 .pred_rep_rf_res_1_3                rand_forest 0.701
## 6 .pred_rep_logreg_resamples_1_112 logistic_reg 0.417
## 7 .pred_rep_lda_resamples_1_07        discrim_linear 0.272
## 8 .pred_rep_logreg_resamples_1_118 logistic_reg 0.203
## 9 .pred_rep_logreg_resamples_1_015 logistic_reg 0.0441
## 10 .pred_rep_lda_resamples_1_04       discrim_linear 0.00112
```

```
gss_preds <-
  gss_test %>%
  dplyr::select(partyid) %>%
  bind_cols(
    predict(
      gss_stack,
      gss_test,
      members = TRUE
    )
  )

colnames(gss_preds) %>%
  map_dfr(
    .f = accuracy,
    truth = partyid,
    data = gss_preds
  ) %>%
  mutate(member = colnames(gss_preds))
```

```
## # A tibble: 13 x 4
##   .metric .estimator .estimate member
##   <chr>    <chr>      <dbl> <chr>
## 1 accuracy binary      1     partyid
## 2 accuracy binary    0.648 .pred_class
## 3 accuracy binary    0.639 .pred_class_logreg_resamples_1_112
```

```
## 4 accuracy binary      0.641 .pred_class_logreg_resamples_1_073
## 5 accuracy binary      0.639 .pred_class_logreg_resamples_1_034
## 6 accuracy binary      0.639 .pred_class_logreg_resamples_1_015
## 7 accuracy binary      0.643 .pred_class_logreg_resamples_1_097
## 8 accuracy binary      0.636 .pred_class_logreg_resamples_1_118
## 9 accuracy binary      0.570 .pred_class_logreg_resamples_1_100
## 10 accuracy binary     0.637 .pred_class_lda_resamples_1_04
## 11 accuracy binary     0.637 .pred_class_lda_resamples_1_07
## 12 accuracy binary     0.64 .pred_class_rf_res_1_2
## 13 accuracy binary     0.652 .pred_class_rf_res_1_3
```

For a base-line, let's fit a simple logistic regression:

```
glm_logreg <- glm(partyid ~ ., data = gss_train, family = "binomial")
probs_logreg <- predict(glm_logreg, gss_test, type = "response")
preds_logreg <- ifelse(probs_logreg >=.5, 1, 0)
confusion_logreg <- table(preds_logreg, gss_test$partyid)
confusion_logreg
```

```
##
## preds_logreg dem rep
##      0 511 213
##      1 316 410
```

```
accuracy_logreg <- 1 - (confusion_logreg[1,2] + confusion_logreg[2,1]) / nrow(gss_test)
accuracy_logreg
```

```
## [1] 0.6351724
```